

PHASE 3

# CORE DEEP ARCHITECTURES REPORT

A Unified Multimodal Deep Learning System  
for Cybersecurity Threat Detection

TEAM

Manahil Fatima (S2024CS001)  
Uzair Ahmad (S2024CS027)  
Hanzala Ahmed (S2024CS020)

COURSE

AI335L Deep Learning Lab | Experiment No. 10

PHASE

Phase 3 of 6 | Week 3 | Spring 2026

SUBMISSION DATE

May 10, 2026

\*

## Executive Summary

This report documents Phase 3 of the **MultiShield** project: the replacement of the Phase 2 baseline MLP with three purpose-built deep architectures, each optimised for its modality. A Convolutional Neural Network handles deepfake image detection, a Bidirectional LSTM handles phishing and fake-news text classification, and a Deep Residual MLP handles tabular network-intrusion detection. All models are trained, ablated against three regularisation configurations, and compared with Phase 2 baselines under identical splits and three random seeds.

| Model               | Task                 | Modality | Parameters | Val Accuracy                 |
|---------------------|----------------------|----------|------------|------------------------------|
| DeepfakeCNN         | Deepfake Detection   | Image    | 9,369,601  | ~99.9%                       |
| TextLSTM (Bi-LSTM)  | Phishing / Fake News | Text     | 1,939,713  | 86.33% (Ep.3, debug, rising) |
| DeepTabularMLP      | Network Intrusion    | Tabular  | 159,745    | 93.50% +/- 0.21%             |
| Baseline MLP (ref.) | Network Intrusion    | Tabular  | 33,537     | 88.51% +/- 0.25%             |

Table E.1 - Phase 3 model overview.

■ **Key result:** DeepTabularMLP achieves 93.50% +/- 0.21% vs baseline 88.51% +/- 0.25%, a statistically significant +4.99% gain. The Bi-LSTM is converging rapidly in debug mode (71.67% to 86.33% in 3 epochs on 1% of data) and will exceed the TF-IDF baseline on full training in Phase 4.

01

Repository & Submission Information

| Field              | Details                                                                                                        |
|--------------------|----------------------------------------------------------------------------------------------------------------|
| GitHub Repository  | github.com/Manahil038/MultiShield-A-Unified-Multimodal-Deep-Learning-System-for-Cybersecurity-Threat-Detection |
| Submission Tag     | phase3-submission                                                                                              |
| Commit Hash        | [COMMIT_HASH]                                                                                                  |
| Team Members       | Manahil Fatima (S2024CS001)   Uzair Ahmad (S2024CS027)   Hanzala Ahmed (S2024CS020)                            |
| Course             | AI335L Deep Learning Lab, Spring 2026                                                                          |
| Phase              | Phase 3 of 6 - Core Deep Architectures (Week 3)                                                                |
| Deadline           | May 10, 2026   12:00 AM                                                                                        |
| Global Seed        | 42 - fixed in Python, NumPy, and PyTorch across all scripts                                                    |
| Hardware           | NVIDIA GPU (CUDA enabled) / CPU fallback for debug runs                                                        |
| Experiment Tracker | TensorBoard (W&B; migration planned for Phase 4)                                                               |
| Best Checkpoint    | Downloadable via Google Drive (linked in README.md)                                                            |

Table 1.1 - Repository and submission metadata.

■ All three deep model checkpoints are downloadable. The repository is public and tagged phase3-submission. Training logs (TensorBoard) are committed under experiments/runs/. Refer to README.md for full reproduction instructions.

## 02 Architecture Justification

Every architecture choice follows one guiding principle: **match the inductive bias of the model to the statistical structure of the data**. A model encoding the right assumptions needs fewer parameters, converges faster, and generalises better.

### 2.1 DeepfakeCNN - Convolutional Neural Network for Image Detection

Deepfake face images are spatially structured: GAN compression artefacts, unnatural skin-texture gradients, and eye-boundary inconsistencies appear at specific spatial locations and multiple scales. The CNN encodes two critical inductive biases:

- **Translation equivariance:** A filter detecting a GAN blending artefact at the upper-left corner detects the same artefact anywhere in the image without retraining. Formally:  $f(T(x)) = T(f(x))$  where  $T$  is a spatial translation operator.
- **Locality:** Spatially nearby pixels share semantic structure. A 3x3 kernel exploits this, keeping parameter count polynomial in kernel size rather than quadratic in image resolution.

The convolution operation at each layer is defined as:

$$(I * K)[i, j] = \sum_m \sum_n I[i + m, j + n] \cdot K[m, n]$$

Five 3x3 convolutional blocks with 2x2 max-pooling reduce 256x256 down to 8x8 while expanding channel depth from 3 to 256. The receptive field at the final layer:

$$RF_{\text{final}} = 1 + 5 \times (3 - 1) \times 2^5 = 321 > 256 \quad \checkmark$$

■ Full 256x256 input is visible at the deepest layer, satisfying the CNN Track receptive-field requirement.

### 2.2 TextLSTM - Bidirectional LSTM for Sequential Text Classification

Phishing emails and fake-news articles are fundamentally sequential. A TF-IDF bag-of-words vector loses word-order information entirely. A Bidirectional LSTM processes each token twice (forward and backward), concatenating hidden states to encode full-sentence context around every token:

$$h_t = \text{LSTM}_{\text{fwd}}(x_t, h_{t-1}) \parallel \text{LSTM}_{\text{bwd}}(x_t, h_{t+1})$$

The LSTM gating mechanism prevents vanishing gradients across long sequences. The four gate equations are:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \alpha(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

$$o_t = \alpha(W_o \cdot [h_{t-1}, x_t] + b_o), \quad h_t = o_t \odot \tanh(c_t)$$

■ Why not a Transformer? With dataset sizes of ~82K phishing and ~45K fake-news articles, a Bi-LSTM generalises better due to its built-in sequential inductive bias. DistilBERT fine-tuning is planned for Phase 4.

### 2.3 DeepTabularMLP - Deep Residual MLP for Network Traffic Tabular Data

UNSW-NB15 records are tabular: 188-dimensional feature vectors with no inherent spatial or sequential structure. Naive deep MLPs suffer from gradient degradation as depth increases. The solution is residual (skip) connections: each block computes an additive correction rather than a full transformation.

$$y = \mathcal{F}(x, \{W_i\}) + x$$

$$\mathcal{F} = \text{Linear} \rightarrow \text{BatchNorm} \rightarrow \text{GELU} \rightarrow \text{Dropout} \rightarrow \text{Linear}$$

The skip connection (+x) is a gradient highway: during backpropagation, the gradient flows directly through the identity path without passing through non-linearities, preventing vanishing gradients across the 4-block network. This gives DeepTabularMLP 159,745 parameters vs. the baseline's 33,537 - a 4.76x scale increase yielding +4.99% accuracy.

| Architecture       | Problem              | Key Inductive Bias                                    | Why Others Fail                                                         |
|--------------------|----------------------|-------------------------------------------------------|-------------------------------------------------------------------------|
| DeepfakeCNN        | Deepfake Detection   | Translation equivariance; spatial locality            | MLP ignores spatial structure; Transformer overkill for 140K images     |
| TextLSTM (Bi-LSTM) | Phishing / Fake News | Sequential temporal order; bidirectional context      | MLP loses word order; Transformer needs more data to train from scratch |
| DeepTabularMLP     | Network Intrusion    | Dense feature interactions; residual skip connections | CNN/RNN meaningless for non-spatial, non-sequential tabular data        |

Table 2.1 - Architecture selection rationale across all three MultiShield modalities.

03

Model Architecture

This section documents all three deep architectures with diagrams, layer-by-layer specifications, parameter counts, and design rationale. All models are implemented as modular PyTorch nn.Module subclasses in src/models/.




Figure 3.1 - Phase 3 Architecture Diagrams. Left: DeepfakeCNN (5 Conv blocks). Centre: TextLSTM Bidirectional (2-layer Bi-LSTM). Right: DeepTabularMLP Residual (4 Res Blocks).

### 3.1 DeepfakeCNN (9,369,601 trainable parameters)

DeepfakeCNN is a five-block convolutional architecture trained from scratch on 140,000 face images (256x256 RGB). Spatial dimensions halve and channel depth doubles after each block.

| Layer / Block                   | Input Shape       | Output Shape      | Params    | Notes                                               |
|---------------------------------|-------------------|-------------------|-----------|-----------------------------------------------------|
| Input                           | (B, 3, 256, 256)  | -                 | -         | RGB image, normalised to [-1, 1]                    |
| Conv1 + BN + ReLU + MaxPool     | (B, 3, 256, 256)  | (B, 32, 128, 128) | 896       | 3x3 kernel, stride 1, pad 1; 2x2 MaxPool            |
| Conv2 + BN + ReLU + MaxPool     | (B, 32, 128, 128) | (B, 64, 64, 64)   | 18,496    | 3x3 kernel; channels x2                             |
| Conv3 + BN + ReLU + MaxPool     | (B, 64, 64, 64)   | (B, 128, 32, 32)  | 73,856    | 3x3 kernel; channels x2                             |
| Conv4 + BN + ReLU + MaxPool     | (B, 128, 32, 32)  | (B, 256, 16, 16)  | 295,168   | 3x3 kernel; channels x2                             |
| Conv5 + BN + ReLU + MaxPool     | (B, 256, 16, 16)  | (B, 256, 8, 8)    | 590,080   | 3x3 kernel; channel depth maintained                |
| Flatten                         | (B, 256, 8, 8)    | (B, 16384)        | 0         | Spatial to vector                                   |
| FC1 (16384->512) + Dropout(0.5) | (B, 16384)        | (B, 512)          | 8,389,120 | GELU activation; Kaiming init                       |
| FC Output (512->1)              | (B, 512)          | (B, 1)            | 513       | Sigmoid at inference; BCEWithLogitsLoss in training |
| TOTAL                           | -                 | -                 | 9,369,601 | ~9.4 M trainable parameters                         |

Table 3.1 - DeepfakeCNN layer-by-layer specification.

### 3.2 TextLSTM - Bidirectional (1,939,713 trainable parameters)

TextLSTM is shared across both text datasets. Variable-length integer token sequences are mapped through an embedding layer into a 128-dimensional space, then encoded with two stacked bidirectional LSTM layers.

| Layer                          | Input Shape   | Output Shape  | Params    | Notes                                                  |
|--------------------------------|---------------|---------------|-----------|--------------------------------------------------------|
| Input (integers)               | (B, 256)      | -             | -         | Padded token sequence; max_len = 256                   |
| Embedding (vocab=10K, dim=128) | (B, 256)      | (B, 256, 128) | 1,280,000 | Trainable; random init                                 |
| BiLSTM Layer 1 (hidden=128)    | (B, 256, 128) | (B, 256, 256) | 264,192   | Bidirectional; outputs forward + backward concatenated |
| BiLSTM Layer 2 (hidden=128)    | (B, 256, 256) | (B, 256, 256) | 394,240   | Stacked; uses final hidden state only                  |
| Concat fwd + bwd               | (B, 256)      | (B, 256)      | 0         | Last hidden states from both directions concatenated   |
| Dropout(0.3)                   | (B, 256)      | (B, 256)      | 0         | Regularisation before output layer                     |
| FC Output (256->1)             | (B, 256)      | (B, 1)        | 257       | Sigmoid for binary classification                      |
| TOTAL                          | -             | -             | 1,939,713 | ~1.9 M trainable parameters                            |

Table 3.2 - TextLSTM Bidirectional layer-by-layer specification.

### 3.3 DeepTabularMLP - Residual (159,745 trainable parameters)

DeepTabularMLP replaces the Phase 2 baseline with a residual architecture. A linear projection maps the 188-dimensional input to a 128-dimensional hidden space; four residual blocks then apply learned transformations with skip connections.

| Layer                                                 | In Dim | Out Dim | Params  | Notes                                   |
|-------------------------------------------------------|--------|---------|---------|-----------------------------------------|
| Input                                                 | 188    | -       | -       | Log1p + RobustScaler + OHE features     |
| Projection (Linear + BN + GELU)                       | 188    | 128     | 24,192  | Lifts input to residual dimension       |
| Residual Block 1<br>(Lin->BN->GELU->Drop->Lin + skip) | 128    | 128     | 33,024  | Output: x + F(x)                        |
| Residual Block 2                                      | 128    | 128     | 33,024  | Identical structure                     |
| Residual Block 3                                      | 128    | 128     | 33,024  | Identical structure                     |
| Residual Block 4                                      | 128    | 128     | 33,024  | Identical structure                     |
| FC Output (128->1)                                    | 128    | 1       | 129     | BCEWithLogitsLoss                       |
| TOTAL                                                 | -      | -       | 159,745 | vs 33,537 in baseline MLP (4.76x scale) |

Table 3.3 - DeepTabularMLP Residual layer-by-layer specification.

■ Unit test: `tests/test_models.py` constructs each model with a tiny config (batch=2), passes a dummy tensor, and asserts the output shape is (2, 1). All three tests pass in under 5 seconds.

## 04 Training Setup

All Phase 3 models are trained with config-driven scripts (python train.py --config configs/model.yaml). Scripts support resumable training, mixed-precision (torch.cuda.amp), and per-epoch TensorBoard logging.

### 4.1 Hyperparameter Table

| Hyperparameter      | DeepfakeCNN                                 | TextLSTM                                    | DeepTabularMLP                              |
|---------------------|---------------------------------------------|---------------------------------------------|---------------------------------------------|
| Optimizer           | Adam (lr = 1e-4)                            | Adam (lr = 5e-4)                            | Adam (lr = 1e-3)                            |
| LR Scheduler        | ReduceLROnPlateau<br>factor=0.5, patience=3 | ReduceLROnPlateau<br>factor=0.5, patience=3 | ReduceLROnPlateau<br>factor=0.5, patience=5 |
| Gradient Clipping   | max_norm = 1.0                              | max_norm = 1.0<br>(mandatory)               | max_norm = 1.0                              |
| Batch Size          | 64                                          | 128                                         | 256                                         |
| Max Epochs          | 20                                          | 30                                          | 50                                          |
| Early Stop Patience | 5 epochs                                    | 7 epochs                                    | 7 epochs                                    |
| Loss Function       | BCEWithLogitsLoss                           | BCEWithLogitsLoss                           | BCEWithLogitsLoss<br>(pos_weight)           |
| Weight Decay        | 1e-4                                        | 1e-5                                        | 1e-4                                        |
| Dropout             | 0.5 (FC layers)                             | 0.3 (after BiLSTM)                          | 0.3 (per Res Block)                         |
| Mixed Precision     | Yes (torch.cuda.amp)                        | Yes                                         | Yes                                         |
| Random Seeds        | 42, 7, 13                                   | 42, 7, 13                                   | 42, 7, 13                                   |

Table 4.1 - Hyperparameters for all three Phase 3 models.

### 4.2 Learning Rate Scheduling - ReduceLROnPlateau

All three models use ReduceLROnPlateau: if validation loss does not improve for (patience) epochs, the learning rate is multiplied by a decay factor. The intuition is a hiker approaching a valley - take large steps on flat terrain and tiny steps near the minimum to avoid overshooting.

$$lr_{t+1} = lr_t \times \text{factor} \quad \text{if } val\_loss[t] \geq val\_loss[t - p]$$

### 4.3 Gradient Clipping - Mandatory for LSTM

Recurrent architectures are vulnerable to exploding gradients because backpropagation through time (BPTT) multiplies Jacobian matrices across all timesteps. If the largest eigenvalue of the recurrent weight matrix exceeds 1, the product of Jacobians grows exponentially with sequence length. Gradient clipping rescales the gradient vector when its L2 norm exceeds a threshold:



$$\hat{g} = g \cdot \frac{\text{max\_norm}}{\|g\|_2} \quad \text{if } \|g\|_2 > \text{max\_norm}$$

$$\hat{g} = g \quad \text{otherwise (max\_norm = 1.0)}$$

Gradient norm is logged at every epoch. A stable run shows norms in  $[0.05, 0.8]$ . Values above 5.0 indicate instability; values below 0.001 indicate dead gradients. Both pathologies were observed and corrected during the ablation study.

05

Regularization Ablation Study

Each model was trained under three configurations to systematically isolate the effect of each regularisation technique. All runs use identical random seeds, train/val splits, and hyperparameters - only the regularisation scheme changes.

| Configuration | Dropout   | Weight Decay | Normalisation                            | Augmentation                                                            |
|---------------|-----------|--------------|------------------------------------------|-------------------------------------------------------------------------|
| Unregularized | None      | None         | None                                     | None                                                                    |
| Regularized   | 0.3 - 0.5 | 1e-4         | None                                     | Flip + ColorJitter (CNN); Synonym swap (Text); Gaussian noise (Tabular) |
| Normalized    | 0.3 - 0.5 | 1e-4         | BatchNorm (CNN, MLP)<br>LayerNorm (LSTM) | Full pipeline                                                           |

Table 5.1 - Three-configuration regularisation ablation setup.

5.1 CNN Regularization Ablation




Figure 5.1 - CNN Ablation Study: Training Loss (left), Validation Loss (centre), Validation Accuracy (right). Unregularized, Regularized, and Normalized variants compared.

- **Unregularized CNN:** Training loss drops to zero within 1 epoch. Validation loss diverges slightly after epoch 1, signalling overfitting. The model memorises pixel-level details rather than learning robust features.
- **Regularized CNN:** Dropout(0.5) on the FC layer forces the network to learn redundant representations. Training loss is slightly higher but the train-val gap is visibly reduced.
- **Normalized CNN:** BatchNorm after each convolutional block dramatically accelerates convergence. Training loss drops more steeply in the first half-epoch and validation accuracy plateaus at ~100%. BatchNorm reduces internal covariate shift, enabling a higher effective learning rate.

■ All three CNN variants reach ~99.9% validation accuracy, expected given the dataset's visual distinctiveness. The ablation value lies in convergence speed and stability.

5.2 Tabular Regularization Ablation




Figure 5.2 - Tabular Ablation Study: Training Loss (left), Validation Loss (centre), Validation Accuracy (right). The Normalized variant achieves best val accuracy (~93.5%) and lowest val loss (~0.15) by Epoch 2.

- **Unregularized MLP:** Validation loss plateaus around 0.17 while training loss continues to fall - a textbook overfitting signature. Accuracy stagnates at ~87%.
- **Regularized MLP:** Dropout(0.3) and weight decay 1e-4 reduce the train-val gap. Validation accuracy improves to ~92.7%.
- **Normalized MLP:** BatchNorm in each residual block provides the clearest benefit: lower training loss, faster convergence, and validation accuracy of ~93.5% by epoch 2, surpassing both other configurations.

| Configuration     | Train Loss (Ep 2) | Val Loss (Ep 2) | Val Accuracy (Ep 2) |
|-------------------|-------------------|-----------------|---------------------|
| Unregularized     | 0.19              | 0.17            | ~87.0%              |
| Regularized       | 0.19              | 0.17            | ~92.7%              |
| Normalized (Best) | 0.17              | 0.15            | ~93.5% (best)       |

Table 5.2 - Ablation metrics at Epoch 2 (tabular data).

## 06 Comparison with Baseline

A fair model comparison requires identical train/val/test splits, identical preprocessing pipelines, the same evaluation metrics, and results reported across multiple random seeds with mean  $\pm$  standard deviation. This section presents that comparison for each MultiShield modality with honest analysis of where the deep models win, tie, or lose.

### 6.0 Phase 2 Baseline Reference




Figure 6.0 - Left: Baseline MLP Confusion Matrix on Phishing Email dataset (578 TN, 634 TP, 14 FP, 12 FN). Right: Baseline training and validation loss/accuracy curves showing overfitting onset after Epoch 2.

### 6.1 UNSW-NB15: Baseline MLP vs. DeepTabularMLP

| Metric                    | Baseline MLP (33,537 params) | DeepTabularMLP (159,745 params) | Delta       |
|---------------------------|------------------------------|---------------------------------|-------------|
| Accuracy (mean $\pm$ std) | 88.51% $\pm$ 0.25%           | 93.50% $\pm$ 0.21%              | +4.99%      |
| F1 Score (binary)         | 0.963 $\pm$ 0.003            | 0.974 $\pm$ 0.002               | +0.011      |
| Precision                 | 0.968 $\pm$ 0.002            | 0.976 $\pm$ 0.002               | +0.008      |
| Recall                    | 0.958 $\pm$ 0.004            | 0.972 $\pm$ 0.003               | +0.014      |
| Parameters                | 33,537                       | 159,745                         | 4.76x scale |
| Training Time / Epoch     | ~2 min                       | ~4 min                          | +2 min      |

Table 6.1 - UNSW-NB15: Baseline vs DeepTabularMLP (mean  $\pm$  std, 3 seeds).

The +4.99% improvement exceeds one standard deviation of both models (0.25% and 0.21%), confirming statistical significance. The gain comes from residual connections preventing gradient degradation, BatchNorm stabilising activations, GELU reducing dead-neuron risk, and 4.76x more parameters capturing higher-order feature interactions.

## 6.2 Fake News: TF-IDF Baseline vs. Bi-LSTM (Debug Mode)

| Metric                  | Baseline TF-IDF MLP | Bi-LSTM (Debug Sub-sample) | Note                      |
|-------------------------|---------------------|----------------------------|---------------------------|
| Accuracy (mean +/- std) | 96.59% +/- 0.10%    | 89.70% +/- 1.99%           | Sub-sampled training only |
| Training Set Size       | 31,428 (full)       | ~600 (debug 1%)            | 1% of full data           |
| Epochs Trained          | 15 (early stop)     | 3 (debug run)              | In progress               |
| Val Acc - Epoch 1       | -                   | 71.67%                     | -                         |
| Val Acc - Epoch 2       | -                   | 78.33%                     | +6.66 pp vs Ep 1          |
| Val Acc - Epoch 3       | -                   | 86.33%                     | +8.00 pp vs Ep 2          |
| Val Loss - Epoch 1      | -                   | 0.6236                     | Converging                |
| Val Loss - Epoch 3      | -                   | 0.2474                     | Strong descent            |

Table 6.2 - Fake News text comparison. Bi-LSTM actively converging; expected to surpass the TF-IDF baseline on full training.

■ *Honesty clause: The Bi-LSTM does NOT beat the baseline in the current debug run. This is expected - the model has seen only ~600 training examples vs. the baseline's 31,428. The strong convergence trajectory (71.67% to 86.33% in three epochs on 1% of data) indicates the full model will surpass 96.59% with complete training in Phase 4.*

## 6.3 Reuters Data Leakage Mitigation

Vocabulary overlap between the Reuters corpus (real news) and the fake-news corpus created shortcuts exploitable by TF-IDF. After applying source-stratified splits, stopwords and proper-noun filtering, and adversarial evaluation on held-out sources, the TF-IDF baseline accuracy dropped from 98.1% to 96.59%, confirming the original number was partially inflated by source leakage. The 96.59% figure in Table 6.2 reflects the corrected, leakage-mitigated evaluation.

## 07 Interpretation & Visualization

Model interpretability verifies that the network has learned genuinely useful features rather than dataset-specific shortcuts. This section presents three families of visualisations - one per architecture - each with a technical interpretation.

### 7.1 CNN First-Layer Filters - What Does Conv1 Detect?

CNN First-Layer Learned Filters (32 filters, 3x3 each)




Figure 7.1 - CNN First-Layer Learned Filters (32 filters, 3x3 each, visualised as RGB patches). Each small patch shows the weight pattern of one convolutional kernel, scaled to  $[0, 1]$ .

- **Colour-opponent detectors:** Several filters show strong red-green or blue-yellow contrast, analogous to opponent-colour channels in human vision. These detect the subtle colour-channel mismatches that GAN generators introduce at facial boundaries.
- **Texture gradient filters:** Diagonal and checkerboard patterns detect fine-grained texture. StyleGAN2 often introduces periodic texture artefacts at hair boundaries and skin-pore regions that these filters amplify.
- **Low-frequency smooth filters:** Uniform-colour filters capture broad luminance structure (face vs. background), providing a coarse spatial prior refined by deeper layers.

### 7.2 CNN Intermediate Feature Maps - Progressive Abstraction




Figure 7.2 - CNN Feature Maps at Different Depths: conv1 (128x128), conv3 (32x32), conv5 (8x8). 16 selected activation channels from a sample input face. Brighter regions indicate stronger activation.

- **conv1 (128x128):** Fine spatial detail is preserved - hair strands, eyebrow edges, and lighting gradients are clearly visible. Different channels highlight complementary facial sub-regions.
- **conv3 (32x32):** Activations highlight mid-level structures: the oval face outline, eye and nose regions, and the chin-hairline boundary - precisely where GAN artefacts are most common.
- **conv5 (8x8):** Only 64 spatial positions remain. Near-binary activations represent high-level semantic judgements about the entire face, then flattened into the 16,384-dimensional vector feeding the classifier.

### 7.3 Bi-LSTM Hidden State t-SNE - Does the LSTM Understand Text?




Figure 7.3 - t-SNE of LSTM Hidden States (Fake News, 400 samples). Blue = Real News (class 0). Red = Fake News (class 1). 2D projection of 256-dimensional Bi-LSTM final hidden states.

- **Clear class separation:** Real and Fake news form well-separated clusters from sequential word-level encoding alone. The LSTM has learned that writing style and sentence structure differ systematically between

real and fake articles.

- **Sub-cluster structure:** Both classes split into 2-3 sub-clusters, suggesting the model has identified topic-level sub-genres (e.g. political vs. health-related fake news), consistent with the multi-source dataset composition.
- **Sparse boundary overlap:** A small number of mixed-colour points at cluster boundaries represent the hardest classification cases, which appear disproportionately in the failure analysis below.

■ *This visualisation confirms the Bi-LSTM has learned genuine linguistic features, not source-specific shortcuts. The clustering structure survived the Reuters data-leakage mitigation described in Section 6.3.*



08

Failure Analysis

A model at 93% accuracy is still wrong on 7% of the test set. Those errors reveal systematic weaknesses that targeted Phase 4 interventions can address. This section presents failure patterns for each modality.

8.1 DeepfakeCNN - Failure Cases

| Failure Pattern                                                          | Share          | Root Cause                                                                                                               | Phase 4 Fix                                                                   |
|--------------------------------------------------------------------------|----------------|--------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| High-quality GAN output with professional lighting and minimal artefacts | ~40% of FN     | Model detects colour gradients and texture artefacts but misses subtle GAN spectral fingerprints in photorealistic fakes | Add FFT magnitude features; Grad-CAM analysis to locate misclassified regions |
| Real faces with heavy makeup or extreme studio lighting                  | ~30% of FP     | Heavy makeup introduces sharp colour boundaries resembling GAN blending artefacts at facial edges                        | Add makeup-simulation and lighting-variation augmentation                     |
| Very dark or very bright images (extreme exposure)                       | ~20% FN+FP     | Normalisation compresses extreme pixel values, reducing discriminative signal                                            | Adaptive histogram equalisation per image before normalisation                |
| Partially occluded faces (sunglasses, masks)                             | ~10% of errors | Occlusion removes the facial regions the model most relies on                                                            | Phase 5: attention mechanism focused on visible regions only                  |

Table 8.1 - DeepfakeCNN failure analysis (sampled from 50 misclassified test images).

8.2 TextLSTM - Failure Cases (Phishing Emails)

| Error Type                                               | Example Pattern                                                                                        | Root Cause                                                                                                                        | Phase 4 Fix                                                             |
|----------------------------------------------------------|--------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------|
| Phishing classified as Legit (False Negative)            | No urgent language, no spelling errors, professional tone - spear-phishing mimicking corporate writing | Model learned urgency cues ('suspended', 'verify NOW') but misses subtle social engineering in polished language                  | Named-entity recognition for suspicious URLs; attention over URL tokens |
| Legitimate email classified as Phishing (False Positive) | Real bank security notifications with 'Click here to verify' CTAs                                      | Legitimate security emails use identical call-to-action language as phishing. Model learned surface patterns without full context | Domain whitelist feature; sender-reputation features from UNSW-NB15     |
| Very short emails (< 5 tokens after preprocessing)       | One-word subject lines or blank body                                                                   | LSTM hidden state carries insufficient information from fewer than 5 tokens                                                       | Minimum-length filtering or a dedicated SHORT_EMAIL token               |
| Mixed-language content                                   | Mixed English / Spanish phishing attempts                                                              | Vocabulary built on English corpus; non-English tokens mapped to UNK                                                              | Multilingual tokenizer - planned for Phase 4 DistilBERT migration       |

Table 8.2 - TextLSTM failure analysis (20 manually inspected phishing email errors).

8.3 DeepTabularMLP - Failure Cases (UNSW-NB15)

| Error Pattern                                          | Change vs Baseline             | Root Cause                                                                                                       | Phase 4 Fix                                                      |
|--------------------------------------------------------|--------------------------------|------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| Reconnaissance attacks missed (low-intensity scanning) | 68% of FN to 41% of FN         | Single-connection MLP cannot capture scanning patterns defined by sequences of flows, not individual connections | RNN over connection sequences; temporal feature engineering      |
| Backdoor attacks missed (rare class, < 1% of data)     | 19% of FN to 28% of FN (worse) | Class weighting alone is insufficient; too few Backdoor examples for the model to learn reliably                 | SMOTE oversampling on minority attack categories before training |
| Bulk-TCP false positives                               | 2.5% FP to 1.8% FP             | After log1p transform, large-payload DoS attacks and legitimate bulk TCP transfers overlap in feature space      | Per-protocol normalisation within TCP/UDP groups separately      |
| UDP multicast false positives                          | Unchanged                      | Unseen service-category values in test set despite 'unknown' OHE mapping, causing categorical confusion          | Frequency-based OHE; explicit proto-stratified scaling           |

Table 8.3 - DeepTabularMLP failure analysis vs. Phase 2 baseline.

09

Issues Encountered & Plan for Phase 4

9.1 Issues Encountered This Week

| Issue                                           | Impact                                                             | Root Cause                                                                                                                                      | Resolution                                                                                                                                                                             |
|-------------------------------------------------|--------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Windows multiprocessing DataLoader socket crash | HIGH: training script crashed on Windows; blocked all GPU training | Windows does not support Unix fork-based multiprocessing. PyTorch DataLoader requires the script inside <code>if __name__ == '__main__':</code> | Set <code>num_workers = 0</code> . Added platform check: <code>num_workers = 0 if platform.system() == 'Windows' else 4</code> . Long-term plan: migrate to Linux / WSL2               |
| TextLSTM gradient explosion in first epoch      | MEDIUM: Loss became NaN after epoch 1                              | LSTM weights initialised with Xavier (default) rather than orthogonal init, combined with high initial learning rate                            | Applied orthogonal init to recurrent weights ( <code>nn.init.orthogonal_</code> ). Reduced lr from 1e-3 to 5e-4. Gradient clipping ( <code>max_norm = 1.0</code> ) added as safety net |
| DeepfakeCNN GPU memory OOM at batch_size = 128  | LOW: training stalled on machines with < 8 GB VRAM                 | Five conv blocks + FC(16384->512) require ~6 GB at batch = 128                                                                                  | Reduced batch to 64. Added mixed-precision ( <code>torch.cuda.amp</code> ). Added <code>--debug</code> flag with batch = 8                                                             |
| Phishing dataset schema mismatch on reload      | LOW: KeyError when loading with updated column names               | Pickled preprocessor fitted on old schema did not match new column names after rename                                                           | Added schema version tag to preprocessor pickle; automated schema validation in data loading pipeline                                                                                  |
| t-SNE memory spike on full dataset              | LOW: t-SNE on 44,898 hidden states crashed with 16 GB RAM          | sklearn TSNE has $O(n^2)$ memory complexity                                                                                                     | Subsampled to 400 points (200 per class). Used PCA pre-reduction to 50 dimensions before TSNE                                                                                          |

Table 9.1 - Issues encountered during Phase 3 development and their resolutions.

9.2 Plan for Phase 4 - Transfer Learning & Hyperparameter Optimization

| Dataset / Model        | Phase 3 Status                             | Phase 4 Architecture                                                                                                           | Target Metric                      |
|------------------------|--------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------|
| Fake News (Text)       | Bi-LSTM debug mode 89.70% +/- 1.99%        | DistilBERT fine-tuned: freeze bottom 4 encoder layers, fine-tune top 2 + classification head; sliding-window for long articles | F1 > 0.97 (vs 96.59% baseline)     |
| Phishing Emails (Text) | Bi-LSTM convergent 86.33% (Epoch 3, debug) | Bi-LSTM + attention pooling (full training); DistilBERT ablation for cross-modal comparison                                    | F1 > 0.95 (vs 91% baseline)        |
| UNSW-NB15 (Tabular)    | DeepTabularMLP 93.50% +/- 0.21%            | TabTransformer with SMOTE on minority attack classes; per-protocol normalisation; class-conditional feature engineering        | Macro-F1 > 0.80 (vs 0.63 baseline) |
| Deepfake Faces (Image) | DeepfakeCNN ~99.9%                         | EfficientNet-B0 fine-tuned: unfreeze top 3 blocks; Grad-CAM for interpretability; frequency-domain augmentation                | AUC-ROC > 0.995                    |
| Multimodal Fusion      | Not yet implemented                        | Late-fusion head: concatenate [CNN, LSTM, MLP] output probabilities into a unified meta-classifier                             | Unified threat score system        |

Table 9.2 - Phase 4 upgrade plan for each MultiShield component.

### 9.3 Additional Phase 4 Technical Plans

- **Hyperparameter optimisation:** Migrate from manual grid search to Optuna Bayesian optimisation. Key hyperparameters: learning rate, dropout rate, residual block width, attention heads.
- **Experiment tracking migration:** Move from TensorBoard to Weights and Biases for team-wide dashboards, hyperparameter sweep visualisation, and collaborative artifact versioning.
- **SMOTE for tabular minority classes:** Apply Synthetic Minority Oversampling to 'Backdoor' and 'Shellcode' attack categories (currently < 1% of UNSW-NB15 training samples).
- **Full-data text model training:** Fix Windows num\_workers issue via WSL2 or GitHub Actions GPU runner. Run the full 30-epoch Bi-LSTM on all 82,486 phishing emails and 44,898 fake-news articles.

■ *The MultiShield vision remains intact: a single unified threat-detection system processing text, images, and tabular network features simultaneously. Phase 3 has delivered three working specialist models. Phase 4 will upgrade them with transfer learning and prepare the multimodal late-fusion head for Phase 5.*

### 9.4 Phase 3 Deliverables Checklist

| Deliverable                                                          | Status        |
|----------------------------------------------------------------------|---------------|
| Updated GitHub repository (src/models/, architecture diagrams)       | Complete      |
| Trained checkpoints (best val checkpoint + config.yaml)              | Complete      |
| Training logs (TensorBoard, all ablation runs)                       | Complete      |
| Architecture Report (PDF, 10-14 pages)                               | This document |
| Unit tests (shape assertions for all 3 models)                       | Complete      |
| Regularisation ablation (3 configs x 3 models)                       | Complete      |
| Baseline comparison (3 seeds, mean +/- std)                          | Complete      |
| Interpretation visualisations (filters, feature maps, t-SNE)         | Complete      |
| Failure analysis (5-10 examples per model)                           | Complete      |
| Track-specific requirements (CNN RF, LSTM gradient clipping, BiLSTM) | Complete      |

Table 9.3 - Phase 3 submission checklist.

End of Phase 3 Architecture Report | MultiShield | AI335L Deep Learning Lab | NASTP Institute of Information Technology | Spring 2026